

ISOM3390 Winter 2021 - Take Home Exam

Lee Jae Yeol 20308109

- Task 1: Web Scraping
- Task 2: Data Wrangling
- Task 3: Data Visualization

Remember to change the `author :` field on this Rmd file to your name and SID.

In assignment 2, we used several datasets provided in the `maps` package and the data contained in **USPresidential08-16.csv** to create a Choropleth map that shows the US 2016 presidential election result at the county level (except for Alaska and Hawaii).

The 2020 presidential election result was called last month. The Democratic candidate, Joe Biden, flipped 5 states (Michigan, Wisconsin, Arizona, Pennsylvania, and Georgia) won by his incumbent counterpart, Donald Trump, in 2016, and won the election. So, it could be interesting to view how changes in county-level result contributed to the former vice president's triumph in the just-passed election.

Task 1: Web Scraping

Many major U.S.-based news websites have published county-level presidential election results (e.g., **www.nbcnews.com**, **www.usatoday**, etc.). So, we can pull the county-level result data from one of these websites to assemble a dataset for later analyses.

Note that, on these websites, the county-level result data is often organized by state and presented by different webpages. And there's often a pattern existing in the URLs of these webpages.

The following gives a list of sample URLs on **www.nbcnews.com** for webpages that present county-level results by state:

```
"https://www.nbcnews.com/politics/2020-elections/alabama-president-results"
"https://www.nbcnews.com/politics/2020-elections/arizona-president-results"
"https://www.nbcnews.com/politics/2020-elections/arkansas-president-results"
"https://www.nbcnews.com/politics/2020-elections/california-president-results"
"https://www.nbcnews.com/politics/2020-elections/colorado-president-results"
"https://www.nbcnews.com/politics/2020-elections/connecticut-president-results"
"https://www.nbcnews.com/politics/2020-elections/delaware-president-results"
"https://www.nbcnews.com/politics/2020-elections/district-of-columbia-president-results"
...
"https://www.nbcnews.com/politics/2020-elections/west-virginia-president-results"
"https://www.nbcnews.com/politics/2020-elections/wisconsin-president-results"
"https://www.nbcnews.com/politics/2020-elections/wyoming-president-results"
```

To form these URLs, you can first use the `state` dataset in `maps` to get the names of all states in the United States mainland (without Alaska and Hawaii), and then use proper `stringr` functions to concatenate them with strings that capture the common pattern (so you don't need to scrape data for Alaska and Hawaii).

```
# Form the URLs for NBC News webpages that present county-level results
# Name the result url
# You are not allowed to use any for/while/repeat loop in this chunk

state <- as_tibble(map_data("state"))
revised_statenames <- str_replace_all(unique(state$region), "( )", "-")

urls <- str_c("https://www.nbcnews.com/politics/2020-elections/", revised_statenames, "-president-results")
head(urls, n = 10)
```

```
## [1] "https://www.nbcnews.com/politics/2020-elections/alabama-president-results"
## [2] "https://www.nbcnews.com/politics/2020-elections/arizona-president-results"
## [3] "https://www.nbcnews.com/politics/2020-elections/arkansas-president-results"
## [4] "https://www.nbcnews.com/politics/2020-elections/california-president-results"
## [5] "https://www.nbcnews.com/politics/2020-elections/colorado-president-results"
## [6] "https://www.nbcnews.com/politics/2020-elections/connecticut-president-results"
## [7] "https://www.nbcnews.com/politics/2020-elections/delaware-president-results"
## [8] "https://www.nbcnews.com/politics/2020-elections/district-of-columbia-president-results"
## [9] "https://www.nbcnews.com/politics/2020-elections/florida-president-results"
## [10] "https://www.nbcnews.com/politics/2020-elections/georgia-president-results"
```

Once you have the URLs ready, start the Web scraping process with the following skeleton code.

Drop the chunk option `eval=F` if you want to knit the `.rmd` file.

```

# Provide appropriate starter code
election_res_2020 <- data.frame(state_county = c(NULL), trump = c(NULL), biden = c(NULL), others = c(NULL))
rd <- rsDriver(chromever = "88.0.4324.27", verbose = F)
remdr <- rd[["client"]]

for (i in seq_along(urls)){

  # Provide your code that mimics needed interactions and extracts and parses HTML source code
  # You are not allowed to use any for/while/repeat loop within this for loop
  # You are only allowed to use *apply functions or leverage R functions' vectorization capabilities
  # whenever an iterative logic needs to be implemented
  # Some states (e.g., Arizona, D.C., etc.) have just a few counties
  # so that there's no "SHOW ALL COUNTIES" button on the corresponding webpages
  # To differentiate them between w/o a "SHOW ALL COUNTIES" button, you can test a condition like
  # length(remdr$findElements('css', ... ))!= 0
  # to check whether the current page contains a "SHOW ALL COUNTIES" button or not
  # You may want to use Sys.sleep(2), which suspends execution of R commands
  # for a specified time interval, e.g., 2 seconds in this example
  # after any browser action so as make sure dynamic contents get fully rendered before extraction
  # Points will be deducted if this loop body contains more than 60 lines (after removing these comments)

  remdr$navigate(urls[i])
  if(length(remdr$findElements(using = "css", ".jsx-1765211304.dib.button-press.founders-cond.lh-none.f4.fw
    6.ttu.clear-blue.bg-white.w-100.pv2"))!= 0) {
    loadmore <- remdr$findElement(using = "css", ".jsx-1765211304.dib.button-press.founders-cond.lh-none.f4.
      fw6.ttu.clear-blue.bg-white.w-100.pv2")
    Sys.sleep(2)
    loadmore$clickElement()
    Sys.sleep(2)
  }
  webpage <- remdr$getPageSource() %>% .[[1]] %>% read_html

  countys <- webpage %>% html_nodes(".publico-txt.truncate.f3.pr2.pr9-m.ls-normal") %>% html_text(trim = TR
    UE) %>% str_replace("100% in$", "") %>% str_to_lower()
  state_county <- str_c(revised_statenames[i],countys, sep = ",")

  votes <- webpage %>% html_nodes(".jsx-3384420229.column-group .jsx-3437879980.dib.number.f2.founders-mono"
    ) %>% html_text(trim = TRUE) %>% str_replace_all(",", "") %>% as.numeric()
  candidates <- webpage %>% html_nodes(".jsx-4189516194.dn.dib-m.mb1") %>% html_text(trim = TRUE)

  len <- length(countys)
  if(candidates[1] == "Donald Trump"){
    election_res_2020 <- rbind(election_res_2020, data.frame(state_county = state_county, trump = c(votes[1:
      len]), biden = c(votes[(len+1):(len*2)]),
      others = c(votes[(len*2+1):(length(votes))]) %>% matrix(nro
        w = len, ncol= length(candidates)-2) %>% rowSums()))
  } else if(candidates[2] == "Donald Trump"){
    election_res_2020 <- rbind(election_res_2020, data.frame(state_county = state_county, biden = c(votes[
      1:len]), trump = c(votes[(len+1):(len*2)]),
      others = c(votes[(len*2+1):(length(votes))]) %>% matrix(nro
        w = len, ncol= length(candidates)-2) %>% rowSums()))
  }
}

```

```

}
}

# Provide your code to organize the scraped data in a tidy data frame named election_res_2020
# Export the data to a same-name csv file (your submission should include this csv file)

election_res_2020 <- as_tibble(election_res_2020)
write.csv(election_res_2020, "election_res_2020.csv")

remdr$close()
rd$server$stop()

```

```
## [1] TRUE
```

```
system("taskkill /im java.exe /f", intern=FALSE, ignore.stdout=FALSE)
```

```
## [1] 0
```

```
election_res_2020
```

```
## # A tibble: 4,593 x 4
##   state_county      trump biden others
##   <chr>          <dbl> <dbl> <dbl>
## 1 alabama,autauga  19838  7503   429
## 2 alabama,baldwin 83544 24578  1557
## 3 alabama,barbour  5622  4816    80
## 4 alabama,bibb    7525  1986    84
## 5 alabama,blount  24711  2640   237
## 6 alabama,bullock  1146  3446    21
## 7 alabama,butler   5458  3965    65
## 8 alabama,calhoun 35101 15216   666
## 9 alabama,chambers 8753  6365   166
## 10 alabama,choke 10583  1624    94
## # ... with 4,583 more rows
```

election_res_2020 that stores the scraped data should look like the following:

```
# A tibble: 4,588 x 4
  state_county      trump biden others
  <chr>          <dbl> <dbl> <dbl>
1 alabama,autauga  19838  7503   429
2 alabama,baldwin 83544 24578  1557
3 alabama,barbour  5622  4816    80
4 alabama,bibb    7525  1986    84
5 alabama,blount  24711  2640   237
6 alabama,bullock  1146  3446    21
7 alabama,butler   5458  3965    65
8 alabama,calhoun 35101 15216   666
9 alabama,chambers 8753  6365   166
10 alabama,choke 10583  1624    94
# ... with 4,578 more rows
```

As you can see, each row corresponds to one county in this data frame. The `trump` and `biden` columns store the numbers of votes received by the 2 major candidates, while the `others` column stores that received by all remaining candidates.

Note: even if you fail to get exactly the same number of rows and the same set of values displayed above, you can still move forward with your imperfect output, or even a fabricated data frame of the same structure (for example, the 2012 presidential election data contained in `USPresidential08-16.csv`) to complete the remaining tasks. ***Full points may still be awarded to sensical answers to the remaining problems.***

Task 2: Data Wrangling

Biden flipped 5 states that Trump won in 2016. They are Michigan, Wisconsin, Arizona, Pennsylvania, and Georgia. The following code creates a character vector for the names of the 5 states:

```
flipped_states <- c('arizona', 'georgia', 'michigan', 'pennsylvania', 'wisconsin')
```

We're going to zoom in on *counties of these 5 states* to inspect changes in vote between the 2016 and 2020 elections.

The 2016 election results of the 5 states can be found in the `USPresidential08-16.csv` file. Use the workaround indicated in assignment 2 to augment observations of interest with the `polynome` column from the `maps::county.fips` dataset.

Rename this column to `state_county`. Name the resulting data frame `election_res_2016`:

```
# Provide your code to create election_res_2016
# You are not allowed to use any for/while/repeat loop in this chunk
# Use inner_join if you want to join two data frames

election_result16 <- read.csv("USPresidential08-16.csv") %>% select(fips = fips_code, ends_with("2016"), -(oth_2016)) %>% inner_join(as.tibble(county.fips), by = "fips") %>% select(ends_with("2016"), state_county = polynome)
```

```
## Warning: `as.tibble()` is deprecated as of tibble 2.0.0.
## Please use `as_tibble()` instead.
## The signature and semantics have changed, see `?as_tibble`.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_warnings()` to see where this warning was generated.
```

```
election_res_2016 <- election_result16 %>% filter(str_extract(election_result16$state_county, "[A-z]+") %in%
  flipped_states) %>% as_tibble()
```

```
election_res_2016
```

```
## # A tibble: 396 x 4
##   total_2016 dem_2016 gop_2016 state_county
##   <int>    <int>    <int> <chr>
## 1    18467     6431    11112 michigan,delta
## 2     6743      904     5676 pennsylvania,fulton
## 3   284832   169169   106559 pennsylvania,delaware
## 4     3572     2695      841 georgia,hancock
## 5     3577     1186     2343 georgia,seminole
## 6    33848    16050    15871 wisconsin,sauk
## 7    96945    34436    58941 pennsylvania,washington
## 8    14757     6774     7239 michigan,leelanau
## 9     6285      619     5561 georgia,brantley
## 10    3486     1156     2158 michigan,baraga
## # ... with 386 more rows
```

The layout of `election_res_2016` should look like the following:

```
# A tibble: 396 x 4
  total_2016 dem_2016 gop_2016 state_county
  <dbl>    <dbl>    <dbl> <chr>
1    18467     6431    11112 michigan,delta
2     6743      904     5676 pennsylvania,fulton
3   284832   169169   106559 pennsylvania,delaware
4     3572     2695      841 georgia,hancock
5     3577     1186     2343 georgia,seminole
6    33848    16050    15871 wisconsin,sauk
7    96945    34436    58941 pennsylvania,washington
8    14757     6774     7239 michigan,leelanau
9     6285      619     5561 georgia,brantley
10    3486     1156     2158 michigan,baraga
# ... with 386 more rows
```

In the 5 flipped states, there are 4 counties whose names are inconsistent between `election_res_2020` and `election_res_2016` :

In `election_res_2020` , their names are "georgia,dekalb" , "michigan,st. clair" , "michigan,st. joseph" , and "wisconsin,st. croix" (they come from NBC News webpages).

while, in `election_res_2016` , they are "georgia,de kalb" , "michigan,st clair" , "michigan,st joseph" , and "wisconsin,st croix" (they come from `maps::county.fips`).

Create a data frame that combines county-level results for the 5 states. Name the resulting data frame

`election_res_1620` :

```
# Provide your code to create election_res_1620
# You are not allowed to use any for/while/repeat loop in this chunk

election_res_2020$state_county[election_res_2020$state_county == "georgia,dekalb"] = "georgia,de kalb"
election_res_2020$state_county[election_res_2020$state_county == "michigan,st. clair"] = "michigan,st clair"
election_res_2020$state_county[election_res_2020$state_county == "michigan,st. joseph"] = "michigan,st joseph"
election_res_2020$state_county[election_res_2020$state_county == "wisconsin,st. croix"] = "wisconsin,st croix"

election_res_1620 <- election_res_2016 %>% inner_join(election_res_2020, by = "state_county") %>% arrange(state_county) %>% mutate(total_2020 = trump + biden + others, dem_2020 = biden, gop_2020 = trump) %>% select(state_county, ends_with("2016"), ends_with("2020")) %>% as_tibble()

election_res_1620
```

```
## # A tibble: 396 x 7
##   state_county    total_2016 dem_2016 gop_2016 total_2020 dem_2020 gop_2020
##   <chr>          <int>    <int>    <int>    <dbl>    <dbl>    <dbl>
## 1 arizona,apache    18659     12196      5315     35183     23293     11442
## 2 arizona,cochise    43147     15291     25036     60473     23732     35557
## 3 arizona,coconino   44929     25308     16573     73346     44698     27052
## 4 arizona,gila       21398      6746     13672     27678      8943     18377
## 5 arizona,graham     11939      3301      8025     14996      4034     10749
## 6 arizona,greenlee    3243       1092      1892      3688       1182      2433
## 7 arizona,la paz      4931       1318      3381       7460       2236      5129
## 8 arizona,maricopa  1201934   549040   590465   2069475  1040774  995665
## 9 arizona,mohave     74189     16485     54656     104705     24831     78535
## 10 arizona,navajo    35409     15362     18165     51783     23383     27657
## # ... with 386 more rows
```

The layout of `election_res_1620` should look like the following:

```
# A tibble: 396 x 7
  state_county    total_2016 dem_2016 gop_2016 total_2020 dem_2020 gop_2020
  <chr>          <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 arizona,apache 18659    12196    5315     35183    23293    11442
2 arizona,cochise 43147    15291    25036    60473    23732    35557
3 arizona,coconino 44929    25308    16573    73346    44698    27052
4 arizona,gila    21398    6746     13672    27678    8943     18377
5 arizona,graham  11939    3301     8025    14996    4034     10749
6 arizona,greenlee 3243     1092     1892     3688     1182     2433
7 arizona,la paz  4931     1318     3381     7460     2236     5129
8 arizona,maricopa 1201934  549040  590465  2069475  1040774  995665
9 arizona,mohave  74189    16485    54656    104705    24831    78535
10 arizona,navajo  35409    15362    18165     51783    23383    27657
# ... with 386 more rows
```

where the column `total_2020` contains the total numbers of votes received at the county level in the 2020 election.

Based on `election_res_1620`, please refer to the **split-apply-combine data analysis workflow** to create a data frame that summarizes the total numbers of votes received by both parties at the state level. Name the resulting data frame `election_res_1620_state`, which should have a layout that looks like the following:

```
# A tibble: 20 x 4
  state    party year    vote
  <chr>    <chr> <chr>    <dbl>
1 Arizona  dem   2016    936250
2 Arizona  dem   2020    1672143
3 Arizona  gop   2016    1021154
4 Arizona  gop   2020    1661686
5 Georgia  dem   2016    1837300
6 Georgia  dem   2020    2473633
7 Georgia  gop   2016    2068623
8 Georgia  gop   2020    2461854
9 Michigan dem   2016    2267373
10 Michigan dem   2020    2804040
11 Michigan gop   2016    2279210
12 Michigan gop   2020    2649852
13 Pennsylvania dem   2016    2844705
14 Pennsylvania dem   2020    3459923
15 Pennsylvania gop   2016    2912941
16 Pennsylvania gop   2020    3378263
17 Wisconsin dem   2016    1382210
18 Wisconsin dem   2020    1630673
19 Wisconsin gop   2016    1409467
20 Wisconsin gop   2020    1610065
```



```
# Provide your code to create election_res_1620_state
# You are not allowed to use any for/while/repeat loop in this chunk

election_res_1620_state <- election_res_1620
name_states <- str_extract(election_res_1620_state$state_county, "[A-z]+")
str_sub(name_states,1,1) <- str_sub(name_states,1,1) %>% str_to_upper()
election_res_1620_state$state_county <- name_states

election_res_1620_state <- election_res_1620_state %>% group_by(state_county) %>% summarise("dem/2016" = sum
  (dem_2016) , "dem/2020" = sum(dem_2020), "gop/2016" = sum(gop_2016), "gop/2020" = sum(gop_2020)) %
>% pivot_longer(cols = -(state_county), names_to = "party", values_to = "vote") %>% separate(col =
  party, into = c("party", "year"))
```

```
## `summarise()` ungrouping output (override with `.groups` argument)
```

```
election_res_1620_state
```

```
## # A tibble: 20 x 4
##   state_county party year    vote
##   <chr>         <chr> <chr>   <dbl>
## 1 Arizona      dem   2016  936250
## 2 Arizona      dem   2020 1672143
## 3 Arizona      gop   2016 1021154
## 4 Arizona      gop   2020 1661686
## 5 Georgia      dem   2016 1837300
## 6 Georgia      dem   2020 2473633
## 7 Georgia      gop   2016 2068623
## 8 Georgia      gop   2020 2461854
## 9 Michigan     dem   2016 2267373
## 10 Michigan    dem   2020 2804040
## 11 Michigan    gop   2016 2279210
## 12 Michigan    gop   2020 2649852
## 13 Pennsylvania dem   2016 2844705
## 14 Pennsylvania dem   2020 3459923
## 15 Pennsylvania gop   2016 2912941
## 16 Pennsylvania gop   2020 3378263
## 17 Wisconsin   dem   2016 1382210
## 18 Wisconsin   dem   2020 1630673
## 19 Wisconsin   gop   2016 1409467
## 20 Wisconsin   gop   2020 1610065
```

Task 3: Data Visualization

With all necessary data ready, first create 2 Choropleth maps for the 2016 and 2020 election results in one figure as shown by the first image posted on Canvas.

As you can see, only the 5 flipped states are colored. And colors represent differences in percentage of votes received by the 2 parties (dem vs. gop).

You need to use geospatial data available in `maps::state` and `maps::county` to draw states' and counties' outlines.

Note that `maps::county` uses the same names for the aforementioned 4 counties as `maps::county.fips`.

Suppose that `p` is the name you to this `ggplot` object. You can directly use the code provided at the end of the following code chunk to style the plot.

Drop the chunk option `eval=F` if you want to knit the `.rmd` file.

```
##>% group_by(state_county) %>% summarise("dem/2016" = sum(dem_2016), "dem/2020" = sum(dem_2020), "gop/2016" = sum(gop_2016), "gop/2020" = sum(gop_2020))

# Provide your code to wrangle the data and create the plot
# You are not allowed to use any for/while/repeat loop in this chunk
# Use inner_join if you want to join two data frames

state <- as_tibble(map_data("state"))
county <- as_tibble(map_data("county"))
county <- county %>% unite(col = state_county, region, subregion, sep = ",")
election_res_1620
```

```
## # A tibble: 396 x 7
##   state_county    total_2016 dem_2016 gop_2016 total_2020 dem_2020 gop_2020
##   <chr>          <int>    <int>   <int>    <dbl>    <dbl>    <dbl>
## 1 arizona,apache    18659    12196    5315     35183    23293    11442
## 2 arizona,cochise   43147    15291   25036     60473    23732    35557
## 3 arizona,coconino  44929    25308   16573     73346    44698    27052
## 4 arizona,gila      21398     6746   13672     27678     8943    18377
## 5 arizona,graham    11939     3301    8025     14996     4034    10749
## 6 arizona,greenlee   3243     1092    1892      3688     1182     2433
## 7 arizona,la paz     4931     1318    3381      7460     2236     5129
## 8 arizona,maricopa 1201934  549040  590465   2069475 1040774  995665
## 9 arizona,mohave    74189    16485   54656    104705    24831    78535
## 10 arizona,navajo   35409    15362   18165     51783    23383    27657
## # ... with 386 more rows
```

```
polygon_data <- election_res_1620 %>% mutate(percentage_change_2016 = (dem_2016 - gop_2016) / total_2016, percentage_change_2020 = (dem_2020 - gop_2020) / total_2020) %>% select(state_county, starts_with("percentage")) %>% pivot_longer(cols = -(state_county), names_to = "year", values_to = "percentage") %>% inner_join(county, by = "state_county")

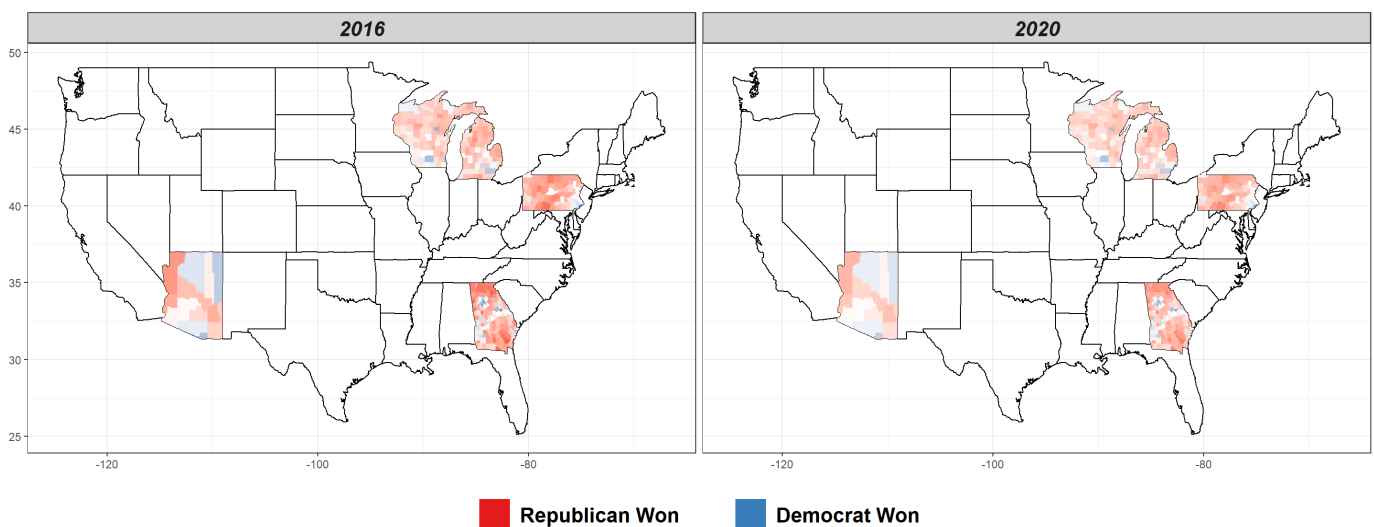
polygon_data$year <- str_replace(polygon_data$year, ".+_ ", "")
polygon_data
```

```
## # A tibble: 21,860 x 7
##   state_county   year percent_change long   lat group order
##   <chr>         <chr>         <dbl> <dbl> <dbl> <dbl> <int>
## 1 arizona,apache 2016           0.369 -109.  36.0   68  3160
## 2 arizona,apache 2016           0.369 -109.  35.0   68  3161
## 3 arizona,apache 2016           0.369 -109.  34.6   68  3162
## 4 arizona,apache 2016           0.369 -109.  33.8   68  3163
## 5 arizona,apache 2016           0.369 -109.  33.8   68  3164
## 6 arizona,apache 2016           0.369 -109.  33.8   68  3165
## 7 arizona,apache 2016           0.369 -109.  33.8   68  3166
## 8 arizona,apache 2016           0.369 -109.  33.8   68  3167
## 9 arizona,apache 2016           0.369 -109.  33.8   68  3168
## 10 arizona,apache 2016           0.369 -109.  33.7   68  3169
## # ... with 21,850 more rows
```

```
p <- ggplot() + geom_polygon(state, mapping = aes(long, lat, group = group), color = "black", fill = NA) +
  geom_polygon(polygon_data, mapping = aes(long, lat, group = group, fill = percent_change)) + facet_wrap(~year)

p + ggtitle("Flipped States: 2016 VS. 2020 Presidential Election Wn") + theme_bw() +
  scale_fill_gradient2(name=NULL, limits = c(-1, 1),
    low = "#e41a1c", high = "#377eb8",
    breaks = c(-1, 1), labels = c("Republican Won", "Democrat Won")) +
  labs(x = NULL, y = NULL) +
  theme(legend.position = "bottom",
    strip.background = element_rect(fill="lightgray", size= 0.8),
    plot.title = element_text(size = 20, face = "bold"),
    strip.text.x = element_text(size = 16, face = "bold.italic"),
    legend.text = element_text(size = 16, face = "bold"),
    legend.spacing.x = unit(0.5, "line"),
    legend.key.size = unit(0.9, "cm")) +
  guides(fill = guide_legend(title.position = "top", title.hjust = 0.5))
```

Flipped States: 2016 VS. 2020 Presidential Election



Please also create 2 bar plots in one figure as shown by the second image posted on Canvas.

Suppose that `h` is the name you to the second `ggplot` object. You can directly use the code provided at the end of the following code chunk to style the plot.

Drop the chunk option `eval=F` if you want to knit the `.rmd` file.

```
# Provide your code to create the plot
# You are not allowed to use any for/while/repeat loop in this chunk

h <- ggplot(election_res_1620_state, aes(x = c(state_county), y = vote, fill = party)) + geom_bar(stat = 'id
  entity', position = "dodge") + facet_wrap(~year)

h + ggtitle("Flipped States: 2016 VS. 2020 Presidential Election Wn") + theme_bw() +
  scale_y_continuous(breaks = c(0, 1000000, 2000000, 3000000),
    labels = c("0", "1,000", "2,000", "3,000")) +
  scale_fill_manual(name=NULL, values = c("dem" = "#377eb8", "gop" = "#e41a1c"),
    labels = c("gop" = "Republican Party", "dem" = "Democrat Party" )) +
  labs(x = NULL, y = "No. of votesWn(in thousands)") +
  theme(legend.position = "bottom",
    strip.background = element_rect(fill="lightgray", size= 0.8),
    plot.title = element_text(size = 20, face = "bold"),
    strip.text.x = element_text(size = 16, face = "bold.italic"),
    legend.text = element_text(size = 16, face = "bold"),
    legend.spacing.x = unit(0.5, "line"),
    legend.key.size = unit(0.9, "cm"),
    axis.title.y = element_text(face="bold.italic", size=18),
    axis.text.x = element_text(size = 14, face="italic"),
    axis.text.y = element_text(size = 14, face="italic")
  ) + guides(fill = guide_legend(title.position = "top", title.hjust = 0.5))
```

Flipped States: 2016 VS. 2020 Presidential Election

